

Introduction to floating point internals and limitations

Dmitrii Milideev, 2020-02-13

Why does it matter?

- We develop our algorithms for real numbers
- Computers can store only a finite subset of reals
- We use float and double types as black boxes, hoping for the best.
- Sometimes we have problems because of that, sometimes not.
 - <https://web.ma.utexas.edu/users/arbogast/misc/disasters.html>
 - Patriot missile crash
 - The short flight of the Ariane 5.
 - Parliamentary elections in Schleswig-Holstein.

What is floating point number

Every real number (except zero) can be represented in the form


$$\pm 1 . d d d . . . \times 2^e$$

Floating point numbers have limited number of digits

$$\pm 1 . d d . . . d \times 2^e$$

IEEE754 representation

Picture taken from https://en.wikipedia.org/wiki/Single-precision_floating-point_format



- Exponent is stored as unsigned integer equal to $\text{<real_exponent>} + 127$
 - This format known as biased representation
 - 127 is constant known exponent bias - every IEEE754 number format has it's own bias.
- Exponent range is -126 to +127. Corresponding biased values are 0x01 and 0xFE.

Special biased exponent value: 0x00

- Representation of zero - all fraction bits are zero
 - Zero can be signed.
- Denormal numbers


$$\pm 0.dd\dots d \times 2^{-127}$$

- Form a fixed-point subset around zero
- Super-slow, don't use them.

Special biased exponent value: 0xFF

- All fraction bits are zero - infinity
 - Positive infinity.
 - Negative infinity.
- Non-zero fraction part - NaN

Overview of available values



Ulp and Ufps

- Let's say we store p fraction bits ($p = 23$ for floats and 52 for doubles)

$$\pm 1 . d_1 d_2 \dots d_p \times 2^e$$

- Actual value of the number

$$x = \pm (2^e \cdot 1 + 2^{e-1} \cdot d_1 + 2^{e-2} \cdot d_2 + \dots + 2^{e-p} \cdot d_p)$$

- $\text{ufp}(x) = 2^e$
 - can be defined for real numbers as $2^{\lceil \log |x| \rceil}$ (according to S.Rump)
- $\text{ulp}(x) = \text{ufp}(x) \cdot 2^{-p}$

Rounding


- In order to put real number into finite representation, it needs to be rounded to a representable number.
- In most cases real number has two neighbors. Rounding can be seen as a process of selection between them.
- Special case if real number is located exactly between two floating-point numbers - need to apply tie-breaking rule.

We have: : 1.10110111111000001111011...

Previous float number : 1.101101111110000011

Next float number : 1.101101111110000100

Median float number : 0.11011



Rounding modes

- **Default** - Round to nearest, ties to even
- Round to nearest, ties away from zero
- Directed modes
 - Towards zero
 - Towards $+\infty$
 - Towards $-\infty$

Rounding error

- When rounding to nearest, absolute error is $\text{ulp}(x) / 2$
- Relative error is bounded by \mathbf{u} - *unit roundoff* (or *machine epsilon*)
- Constant \mathbf{u} is dependent only on a number of digits we are storing.
- \mathbf{u} equals
 - 2^{-24} for float
 - 2^{-53} for double

Rounding happens after every operation

- IEEE754 requires every operation to be exactly rounded
 - Inputs are treated as exact numbers
 - Operations performed with infinite mantissa length
 - Rounded afterwards
- Rounding happens implicitly after *every* arithmetic operation.
 - You get $\circ(a \cdot b)$ instead of $a \cdot b$
- Sometimes you're lucky and result of operation is *exact*
 - Happens when mantissa of the result is short enough to fit your datatype
- Example: for a and b , the rounded multiplication result is somewhere here

$$[(a \cdot b) - \text{ulp}(a \cdot b)/2 ; (a \cdot b) + \text{ulp}(a \cdot b)/2]$$

$$[(a \cdot b) - u * x ; (a \cdot b) + u * x]$$

Examples of exact operations

- All divisions and multiplications by 2^x
 - You modify only exponent not mantissa. Nothing to round
- All integers that fit mantissa have exact representation
 - Multiplications, additions and subtractions are all exact
- All fixed-point numbers operations are exact (e.g. only 4 digits after radix point)
 - Essentially, they are integers with shifted exponent value
 - Additions are exact
 - Multiplication requires $n+m$ digits, if it fits mantissa - result is exact.

Examples of inexact operations

- All money-related data
 - In binary representation, constant 0.1 is $0.0001100110011(0011)$
 - All dollars-and-cents values cannot be stored exactly. They are always rounded (up or down).
 - You always get slightly imprecise results.
- Obviously, results of most divisions and trig functions.

Fused multiply–add (FMA) operations

- $\text{FMA}(a, b, c)$ computes $a \cdot b + c$ with single rounding operation
 - Have implementation in silicon
- You get $\text{○}(a \cdot b + c)$ instead of $\text{○}(\text{○}(a \cdot b) + c)$
- Computations are faster and more accurate with FMA.
 - For example, Eigen always trying to use FMA instructions if possible
- Compilers are pretty conservative with the usage of FMA instructions
 - You have to ask them specifically
- Available since C++11 with `std::fma`

Example: determinant of the 2x2 matrix

We want to compute

$$\det(A) = a_{11} \cdot a_{22} - a_{12} \cdot a_{21}$$

Our data type is 32-bit float from C++.

```
float a11 = 6.0f - 56 * 0x1p-12;
```

```
float a12 = 6.0f - 55 * 0x1p-12;
```

```
float a21 = 18.0f - 57 * 0x1p-12;
```

```
float a22 = 18.0f - 54 * 0x1p-12;
```


Example: determinant of the 2x2 matrix

We want to compute

$$\det(A) = a_{11} \cdot a_{22} - a_{12} \cdot a_{21}$$

Binary representation.

$$\text{float } a_{11} = +1.0111 \ 1111 \ 0010 \ 0000 \ 0000 \ 000 \ * \ 2^2$$

$$\text{float } a_{12} = +1.0111 \ 1111 \ 0010 \ 0100 \ 0000 \ 000 \ * \ 2^2$$

$$\text{float } a_{21} = +1.0001 \ 1111 \ 1100 \ 0111 \ 0000 \ 000 \ * \ 2^4$$

$$\text{float } a_{22} = +1.0001 \ 1111 \ 1100 \ 1010 \ 0000 \ 000 \ * \ 2^4$$

Example: determinant of the 2x2 matrix

Let's compute exact result $a_{11} \cdot a_{22}$

$$a_{11} \cdot a_{22} = +1.1010 \ 1110 \ 1011 \ 0011 \ 0010 \ 1111 \ 01 * 2^6$$

Let's compute rounded result $\circ(a_{11} \cdot a_{22})$

$$\circ(a_{11} \cdot a_{22}) = +1.1010 \ 1110 \ 1011 \ 0011 \ 0011 \ 000 * 2^6$$

Result of multiplication was rounded up.

Example: determinant of the 2x2 matrix

Let's compute exact result $a_{12} \cdot a_{21}$

$$a_{12} \cdot a_{21} = +1.1010 \ 1110 \ 1011 \ 0011 \ 0011 \ 000\mathbf{0} \ \mathbf{1111} \ \mathbf{11} * 2^6$$

Let's compute rounded result $\circ(a_{12} \cdot a_{21})$

$$\circ(a_{12} \cdot a_{21}) = +1.1010 \ 1110 \ 1011 \ 0011 \ 0011 \ 000 * 2^6$$

Result of multiplication was rounded down.

Example: determinant of the 2x2 matrix

Let's compute

$$\circ(a_{11} \cdot a_{22}) - \circ(a_{12} \cdot a_{21})$$

Result is zero.

Correct result is:

$$111 * 0x1p-24$$

$$-6.61612e-06$$

$$-1.1011 \ 1100 \ 0000 \ 0000 \ 0000 \ 000 * 2^{-18}$$

Problems with error accumulation

- Roundoff error in intermediate results might lead to a pretty big error in final result.
 - Most dangerous part is if you have an error accumulation inside a loop.
- Even if your data does not require large mantissa, your intermediate results usually do.
 - For example, comparing x-coordinates of intersection points of two segments might require 5x mantissa length.

Easy error estimation

- Simplest way to estimate roundoff error is mid-point interval arithmetic or ball arithmetic.
 - See: S.M. Rump. Fast and parallel interval arithmetic, 1999.
- Every number is viewed as a random value from eps-neighborhood of some real value:

$$[x - \varepsilon; x + \varepsilon] =: \langle x, \varepsilon \rangle$$

- Arithmetic rules are defined like

$$\langle x, \varepsilon_1 \rangle + \langle y, \varepsilon_2 \rangle = \langle x+y, \varepsilon_1+\varepsilon_2 \rangle$$

$$\langle x, \varepsilon_1 \rangle * \langle y, \varepsilon_2 \rangle = \langle x*y, \varepsilon_1*|y| + \varepsilon_2*|x| + \varepsilon_1*\varepsilon_2 \rangle$$

$$\circ(\langle x, \varepsilon \rangle) = \langle x, \varepsilon + u*(|x| + \varepsilon) \rangle$$

Easy error estimation: example

$$\det(A) = a_{11} \cdot a_{22} - a_{12} \cdot a_{21}$$

- All elements are bounded by some constant $|a_{ij}| < A$
- Treat inputs as exact: a_{ij} estimation is $\langle A, 0 \rangle$
- $a_{ij} \cdot a_{ij}$ estimation is $\langle A \cdot A, 0 \rangle$
- $\circ(a_{ij} \cdot a_{ij})$ estimation is $\langle A^2, \mathbf{u} \cdot A^2 \rangle$
- $\circ(a_{11} \cdot a_{22}) - \circ(a_{12} \cdot a_{21})$ estimation is $\langle 2 \cdot A^2, 2 \cdot \mathbf{u} \cdot A^2 \rangle$
- $\det(A) = \circ(\circ(a_{11} \cdot a_{22}) - \circ(a_{12} \cdot a_{21}))$ estimation is

$$\langle 2 \cdot A^2, 3 \cdot \mathbf{u} \cdot A^2 + 2 \cdot \mathbf{u}^2 \cdot A^2 \rangle$$

Easy error estimation: example

$$\langle A^2, 3 \cdot \mathbf{u} \cdot A^2 + 2 \cdot \mathbf{u}^2 \cdot A^2 \rangle$$

- For float, $\mathbf{u}=2^{-24}$
- Let's say $A=20$
- Error estimation value is

7.152557657263969e-05

Libraries

- GNU GMP
 - <https://gmplib.org/>
- boost::multiprecision
 - https://www.boost.org/doc/libs/1_72_0/libs/multiprecision/doc/html/index.html
- Intel C++ Math Library

Thanks!

- Questions?
- Some good references on exactness approach
 - S. Rump is my favorite author, pretty much everything is from him
<http://www.ti3.tu-harburg.de/rump/>
 - S.M. Rump, T. Ogita, and S. Oishi. Accurate floating-point summation part I: Faithful rounding. <http://www.ti3.tu-harburg.de/paper/rump/RuOgOi07I.pdf>
 - S.M. Rump. Fast and parallel interval arithmetic
<http://www.ti3.tu-harburg.de/paper/rump/Ru99b.pdf>
 - M. Lange and S.M. Rump. Faithfully Rounded Floating-point Computations
<http://www.ti3.tu-harburg.de/paper/rump/LaRu2017b.pdf>
 - K. Ozaki, T. Ogita, S. Oishi. A robust algorithm for geometric predicate by error-free determinant transformation <https://www.sciencedirect.com/science/article/pii/S0890540112000752>
- Big book: J.-M. Muller, Handbook of Floating-Point Arithmetic (not freely available)
<https://www.springer.com/us/book/9783319765259>